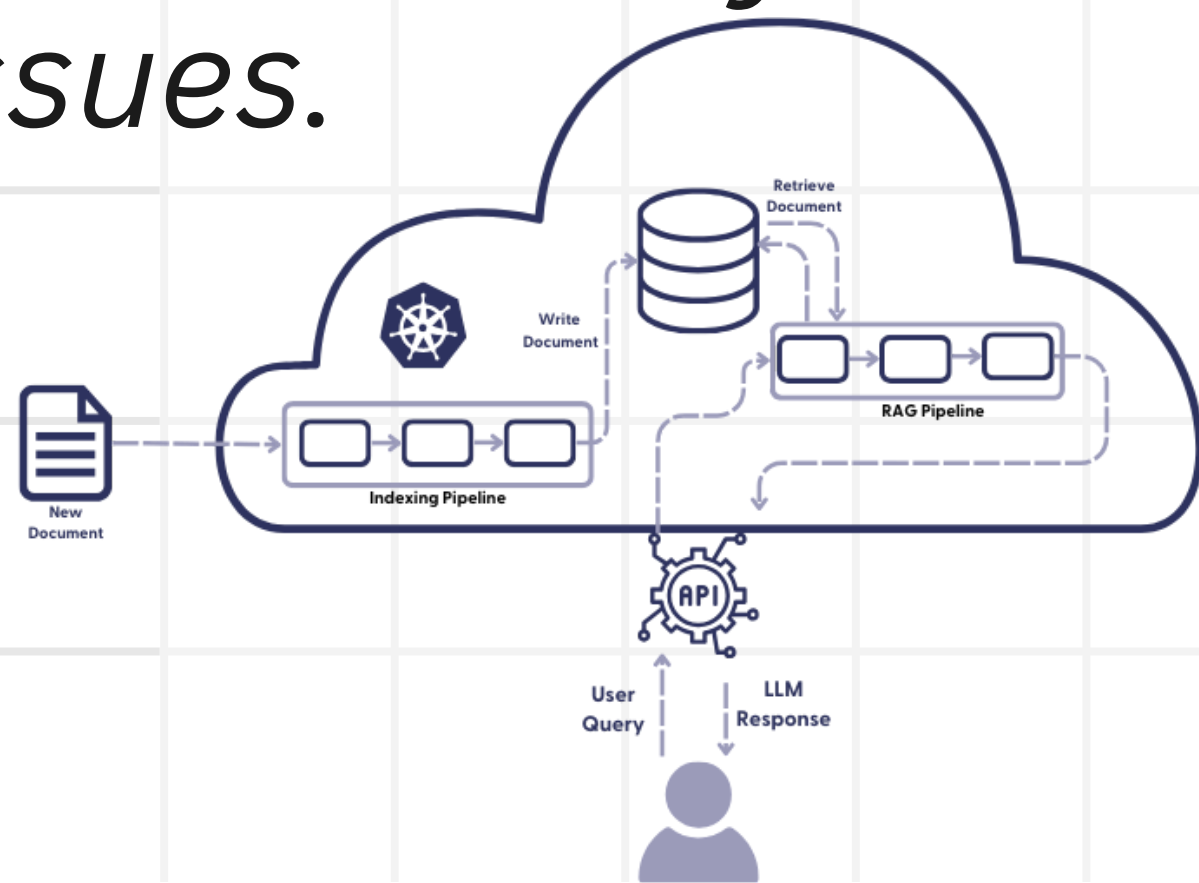


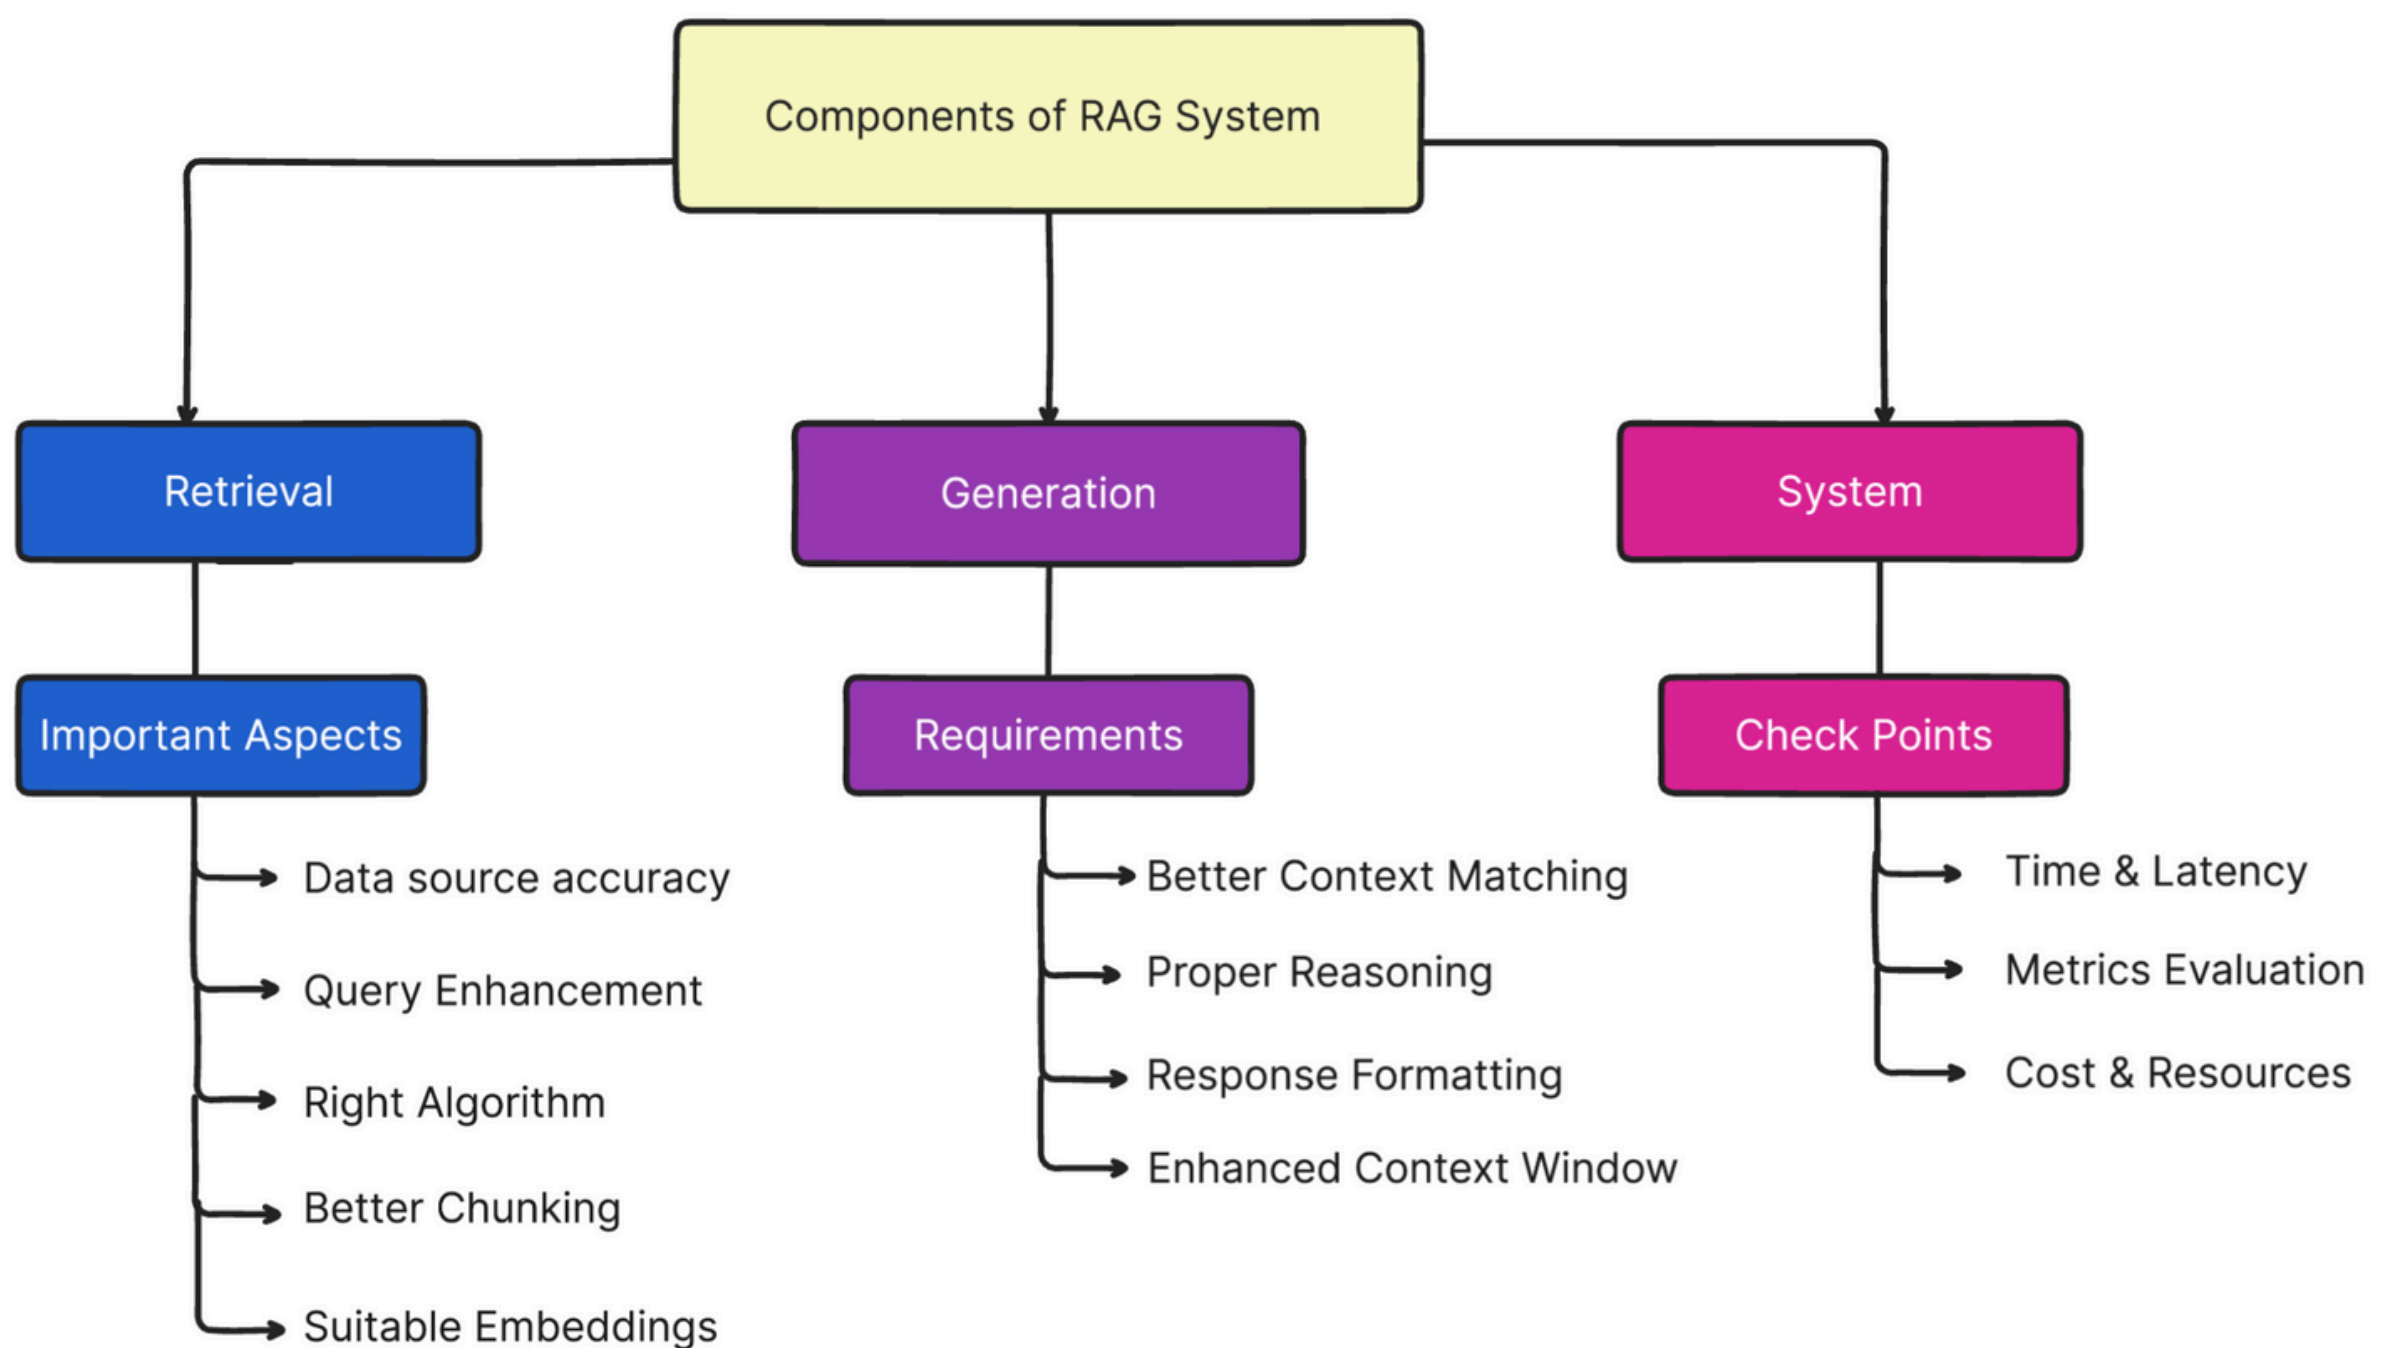


A GUIDE TO BUILDING **A RAG SYSTEM THAT WORKS!**

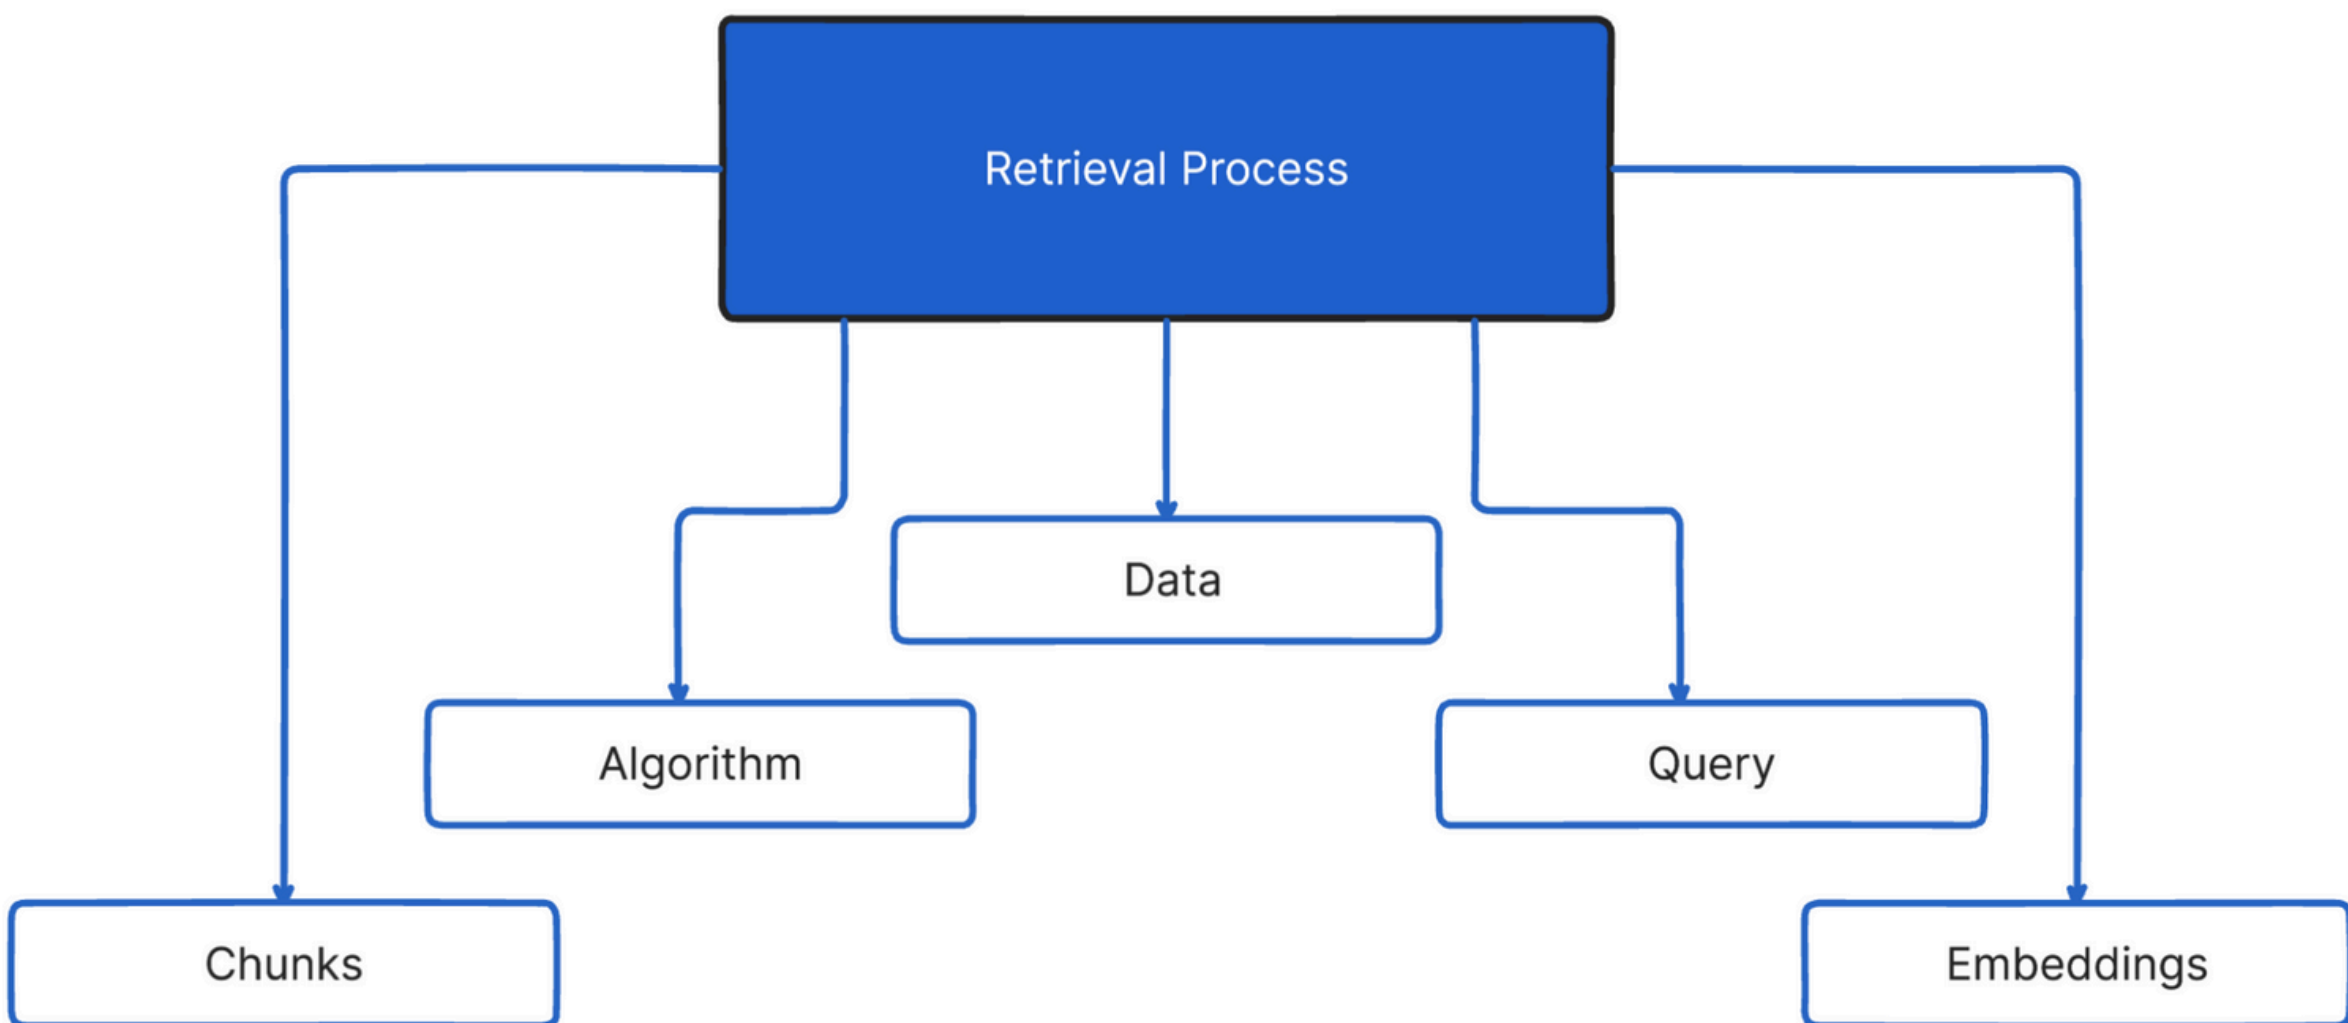
*+ Find solutions to fix all your
RAG system issues.*



COMPONENTS OF A RAG SYSTEM



RETRIEVAL PROCESS



Challenges:

- Data & Query Mismatch
- Search/ Retrieval Algorithm Shortcomings
- Challenges in Chunking
- Embedding Problems

DATA & QUERY

Data & Query Mismatch

Query Ambiguity & Lack of Context

Working with Inaccurate Data Sources

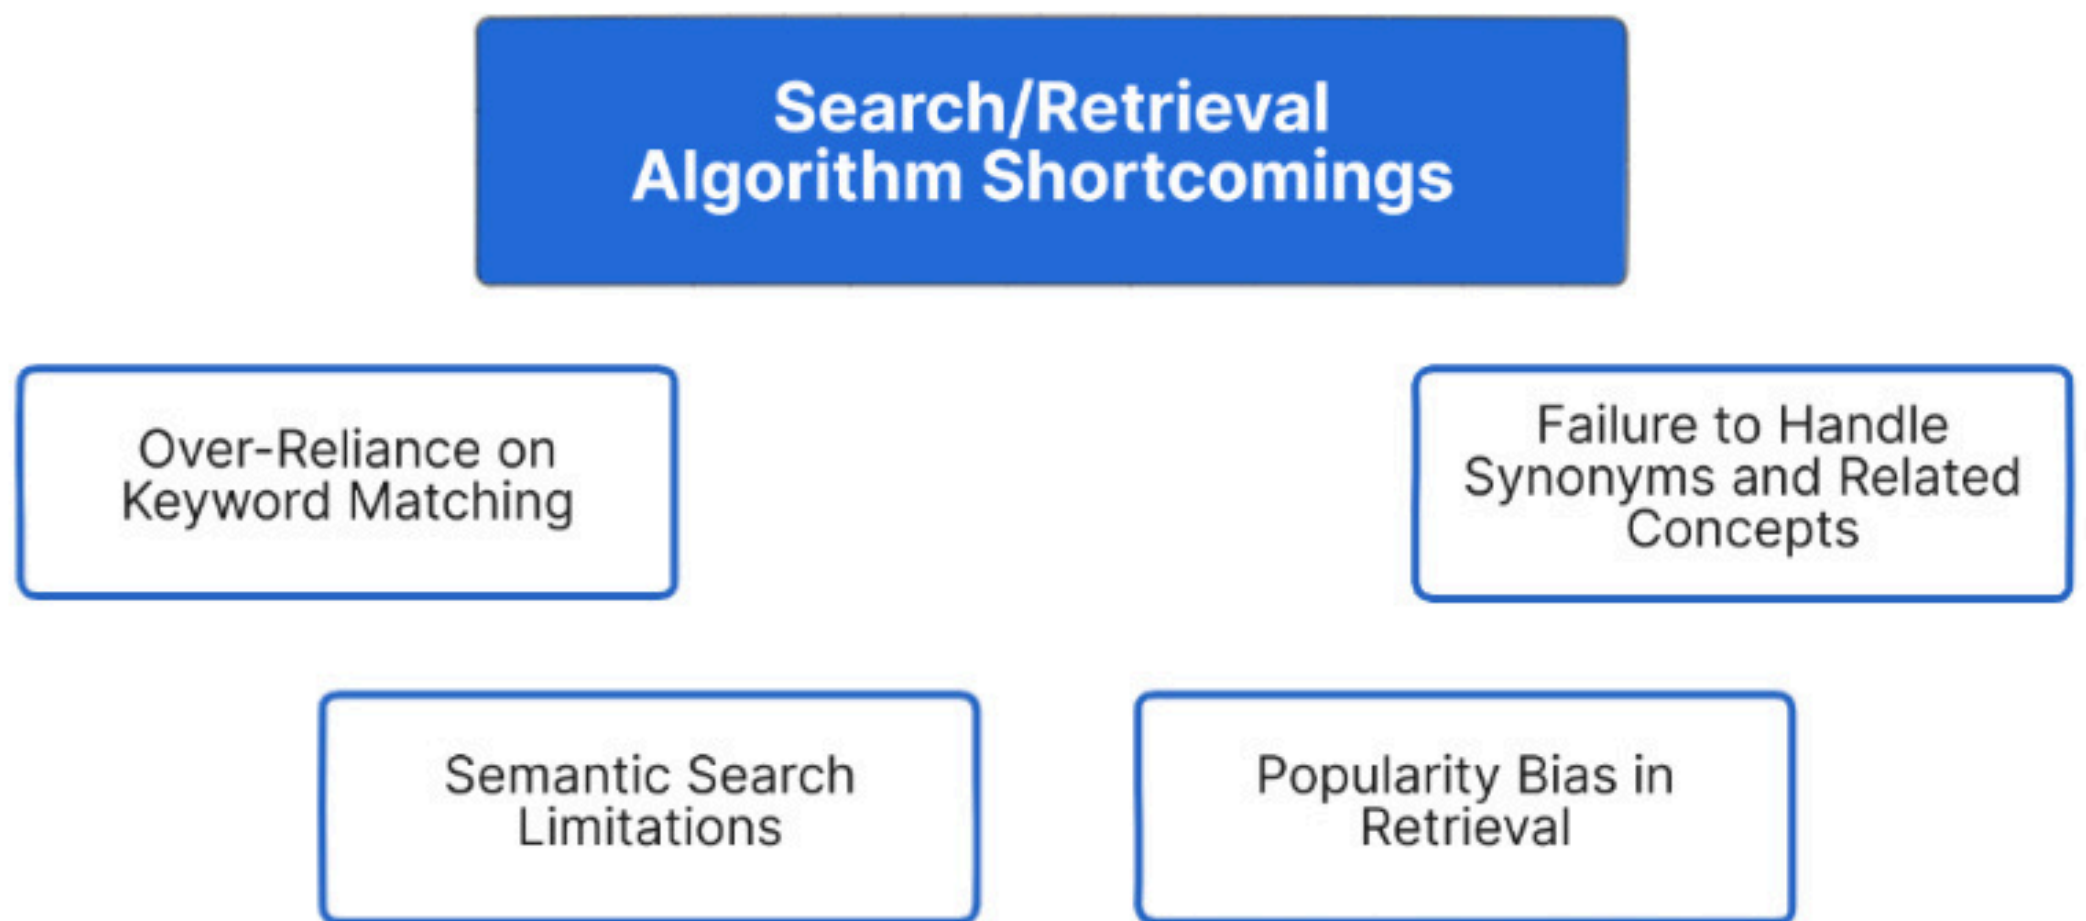
Difficulty with Complex, Multi-faceted Queries

Over-reliance on Keyword Matching

Solution:

- Add Possible Solutions Along with the Query.
- Add Other Similar Queries
- Personalise each query with context
- Consider which data source(s) will be the most relevant for that RAG system.

SEARCH/RETRIEVAL ALGORITHMS



Solution:

- Combine keyword (BM25) and semantic search for balanced results.
- Enhance queries (synonyms, context, rephrasing) for better retrieval.
- Use multiple methods (lexical, dense) with re-ranking for improved coverage and relevance.

CHUNKING

Challenges in Chunking

Inappropriate Chunk
Sizes (Too Large or Too
Small)

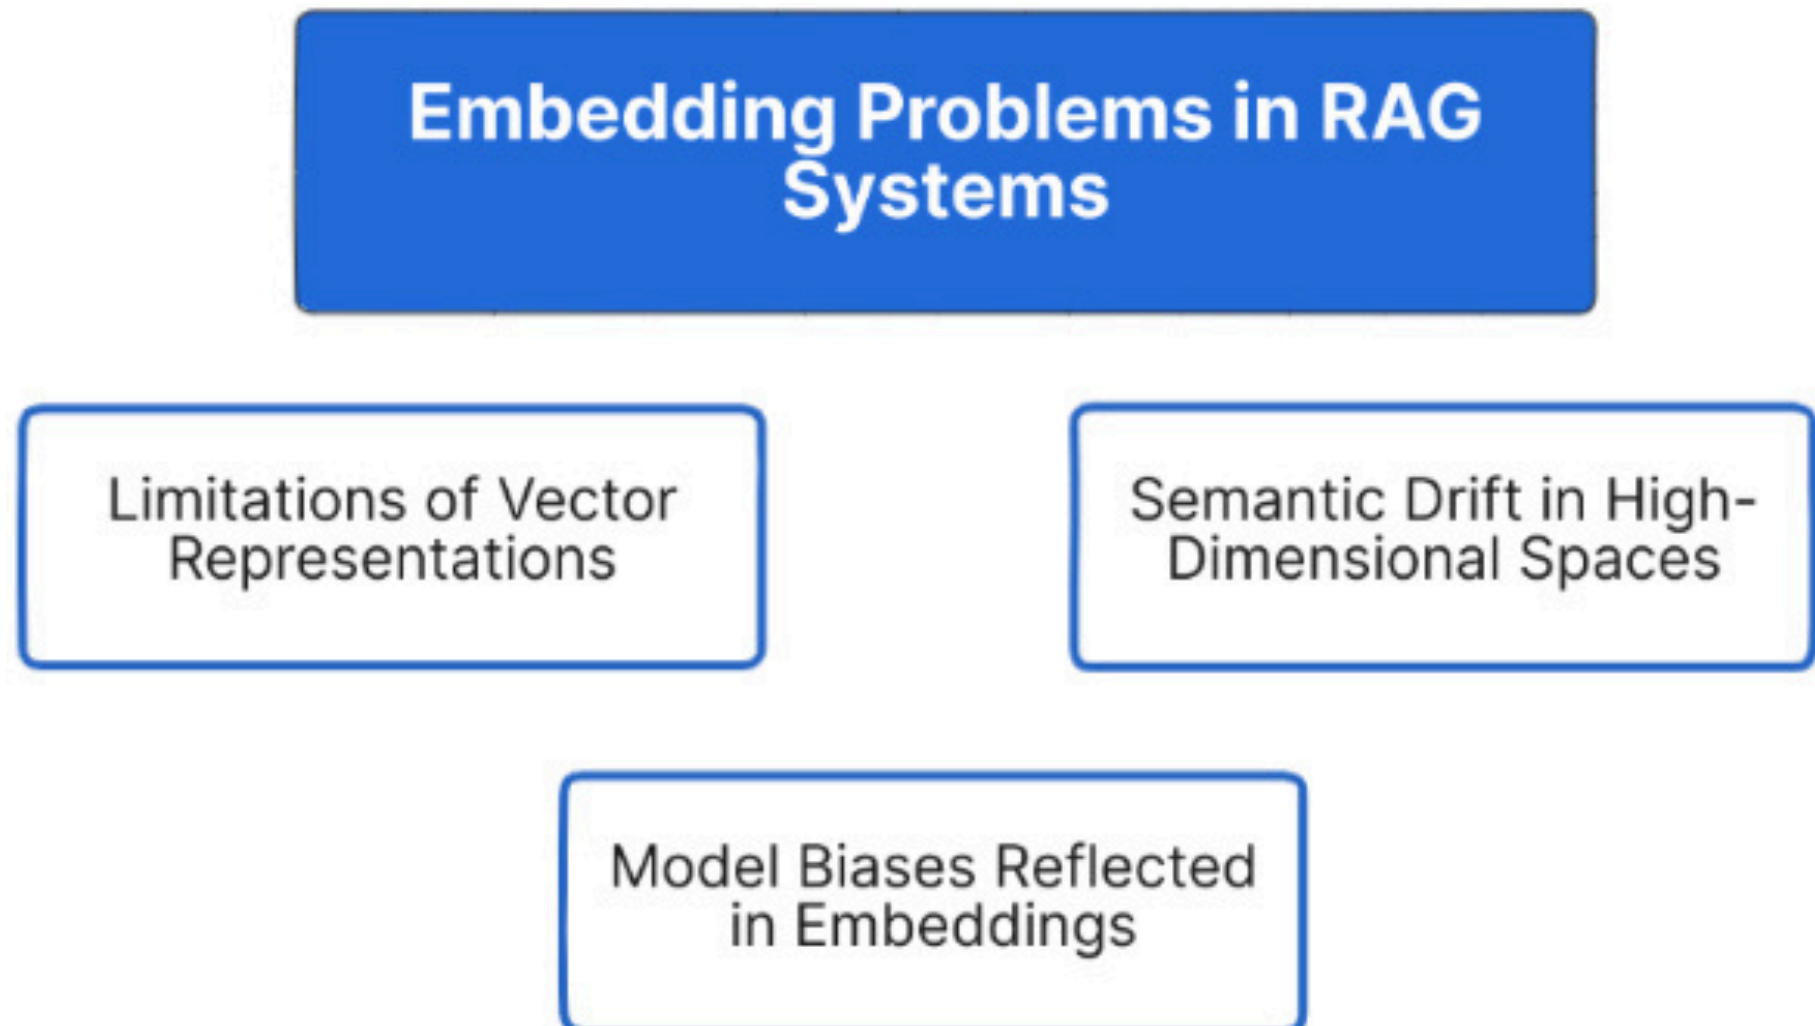
Loss of Context When
Splitting Documents

Failure to Maintain
Semantic Coherence
Across Chunks

Solution:

- Use NLP to find natural breakpoints, creating meaningful chunks.
- Divide structured documents along existing sections/titles.
- Add overlapping text between chunks to maintain context/references.
- Employ AI to adapt chunk size based on topic shifts for relevance.

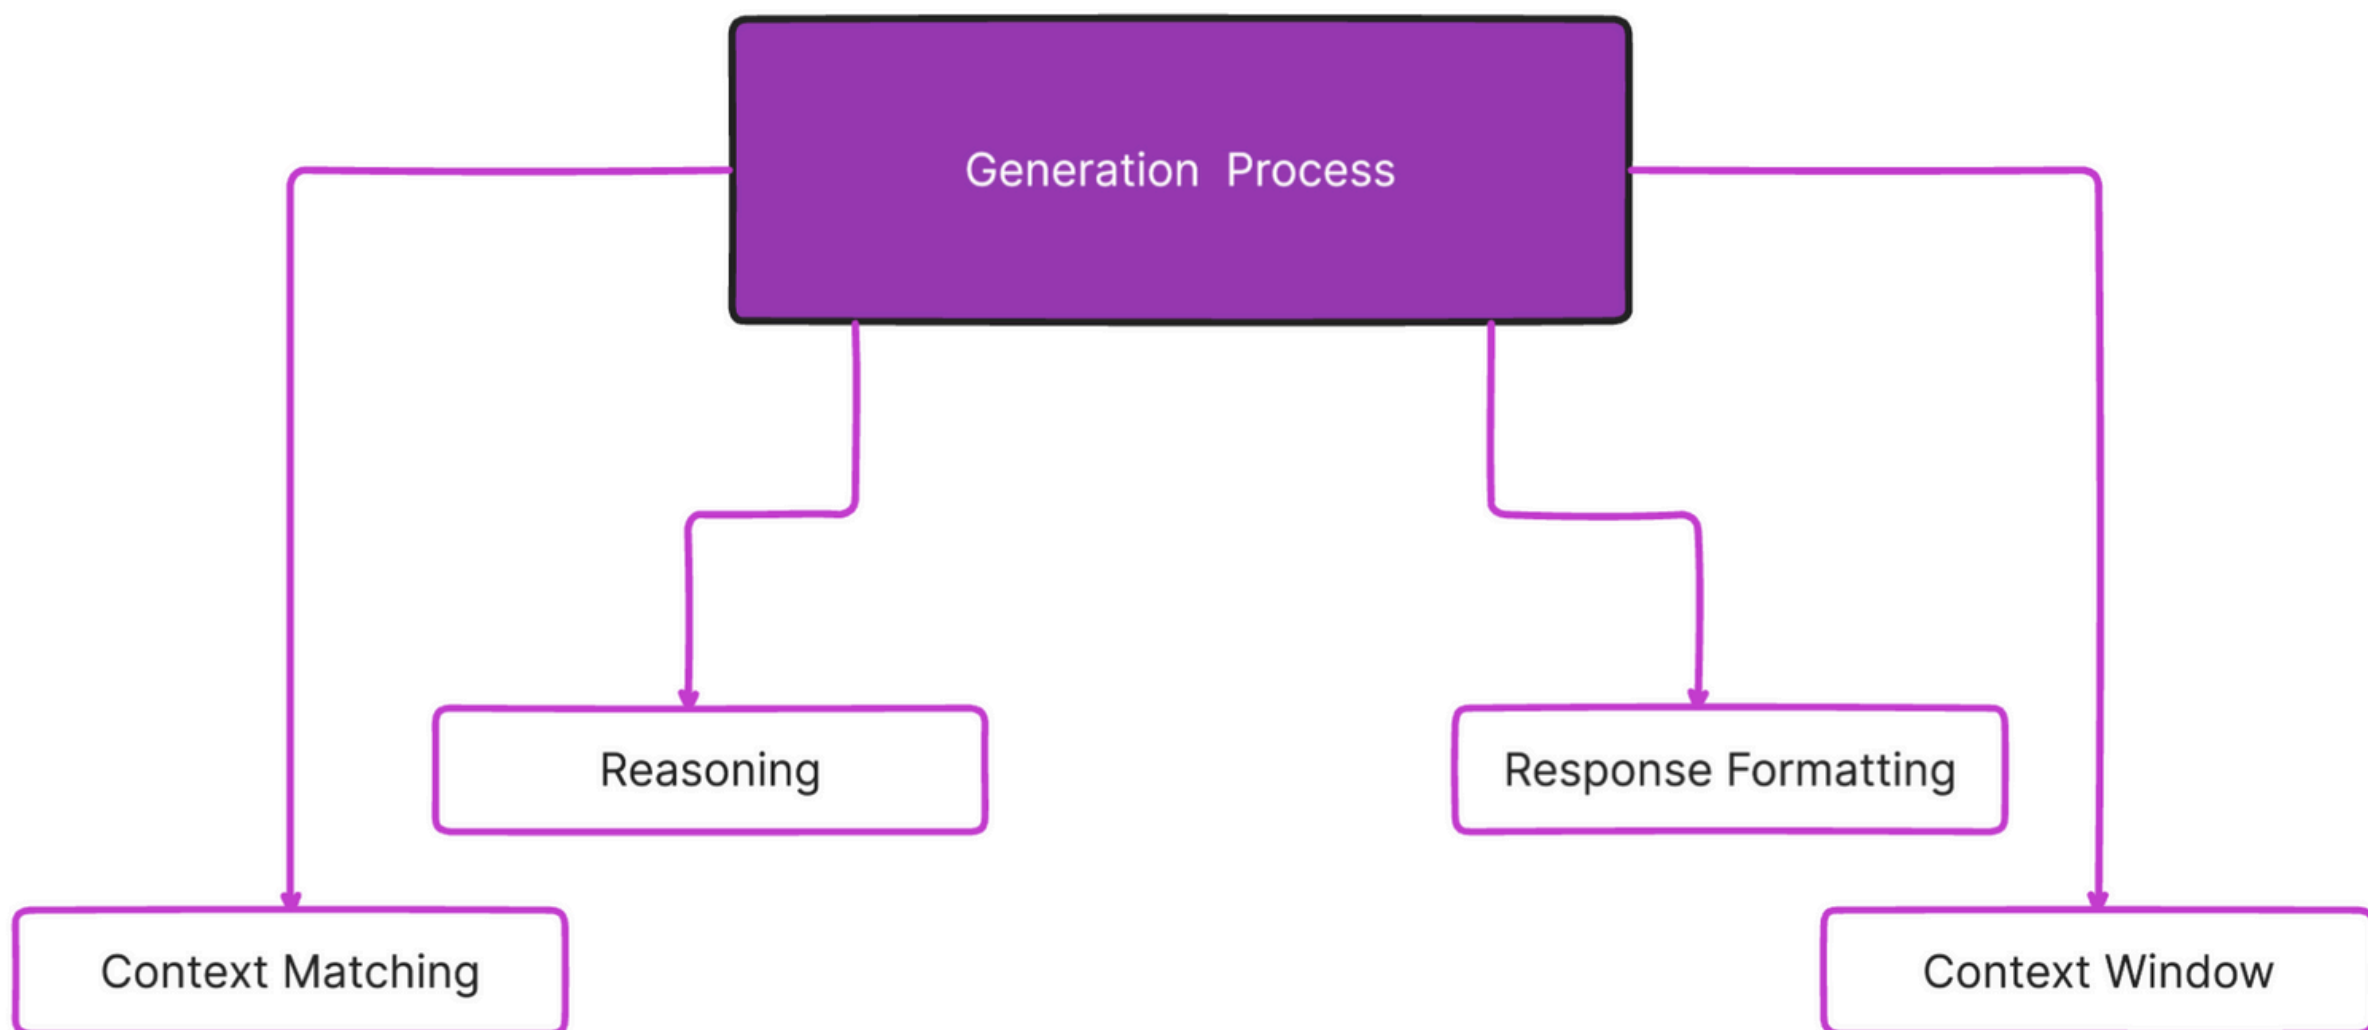
EMBEDDING



Solution:

- Adapt embeddings using domain-specific data for accuracy.
- Re-embed knowledge frequently to stay current.
- Combine traditional and contextual models for better understanding.

GENERATION PROCESS FAILURES



Reasons:

- Context Integration Problems
- Reasoning Limitation
- Response Formatting Issues
- Context Window Utilization

CONTEXT INTEGRATION

Context Integration Problems

Failure to Properly
Incorporate Retrieved
Information

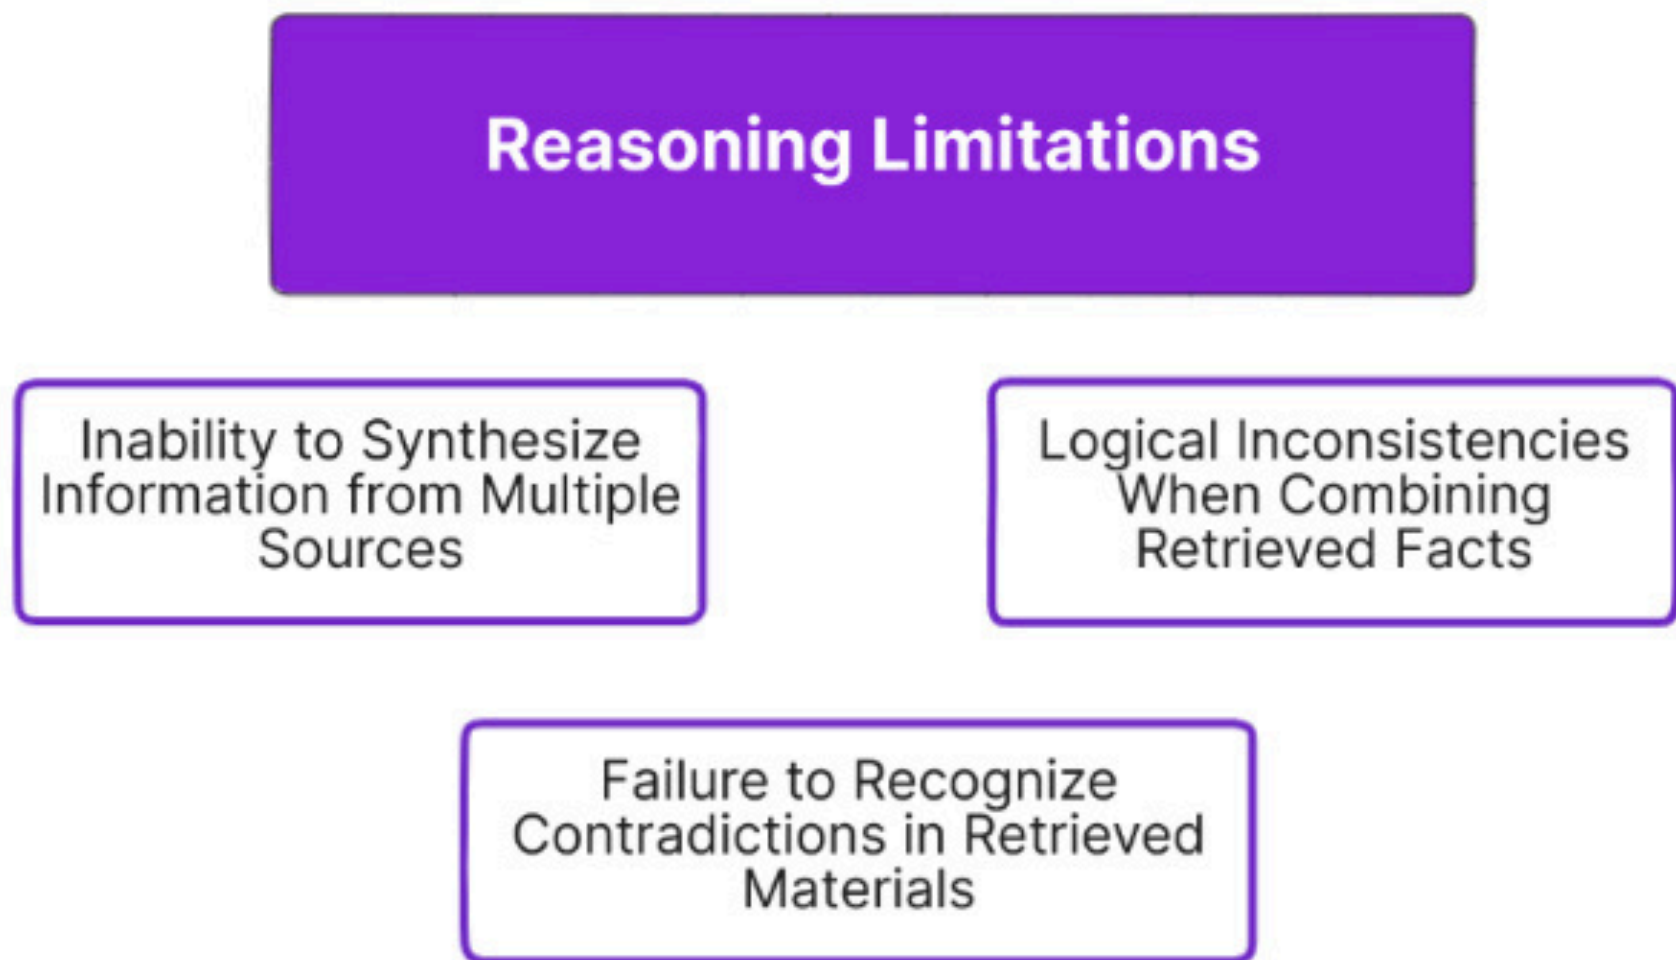
Hallucinations Despite
Having Correct
Information in Context

Over-Reliance on Model's
Parametric Knowledge vs.
Retrieved Information

Solution:

- Supervised FineTuning for Better Grounding
- Fact Verification Post-Processing
- Retrieval-Aware Training

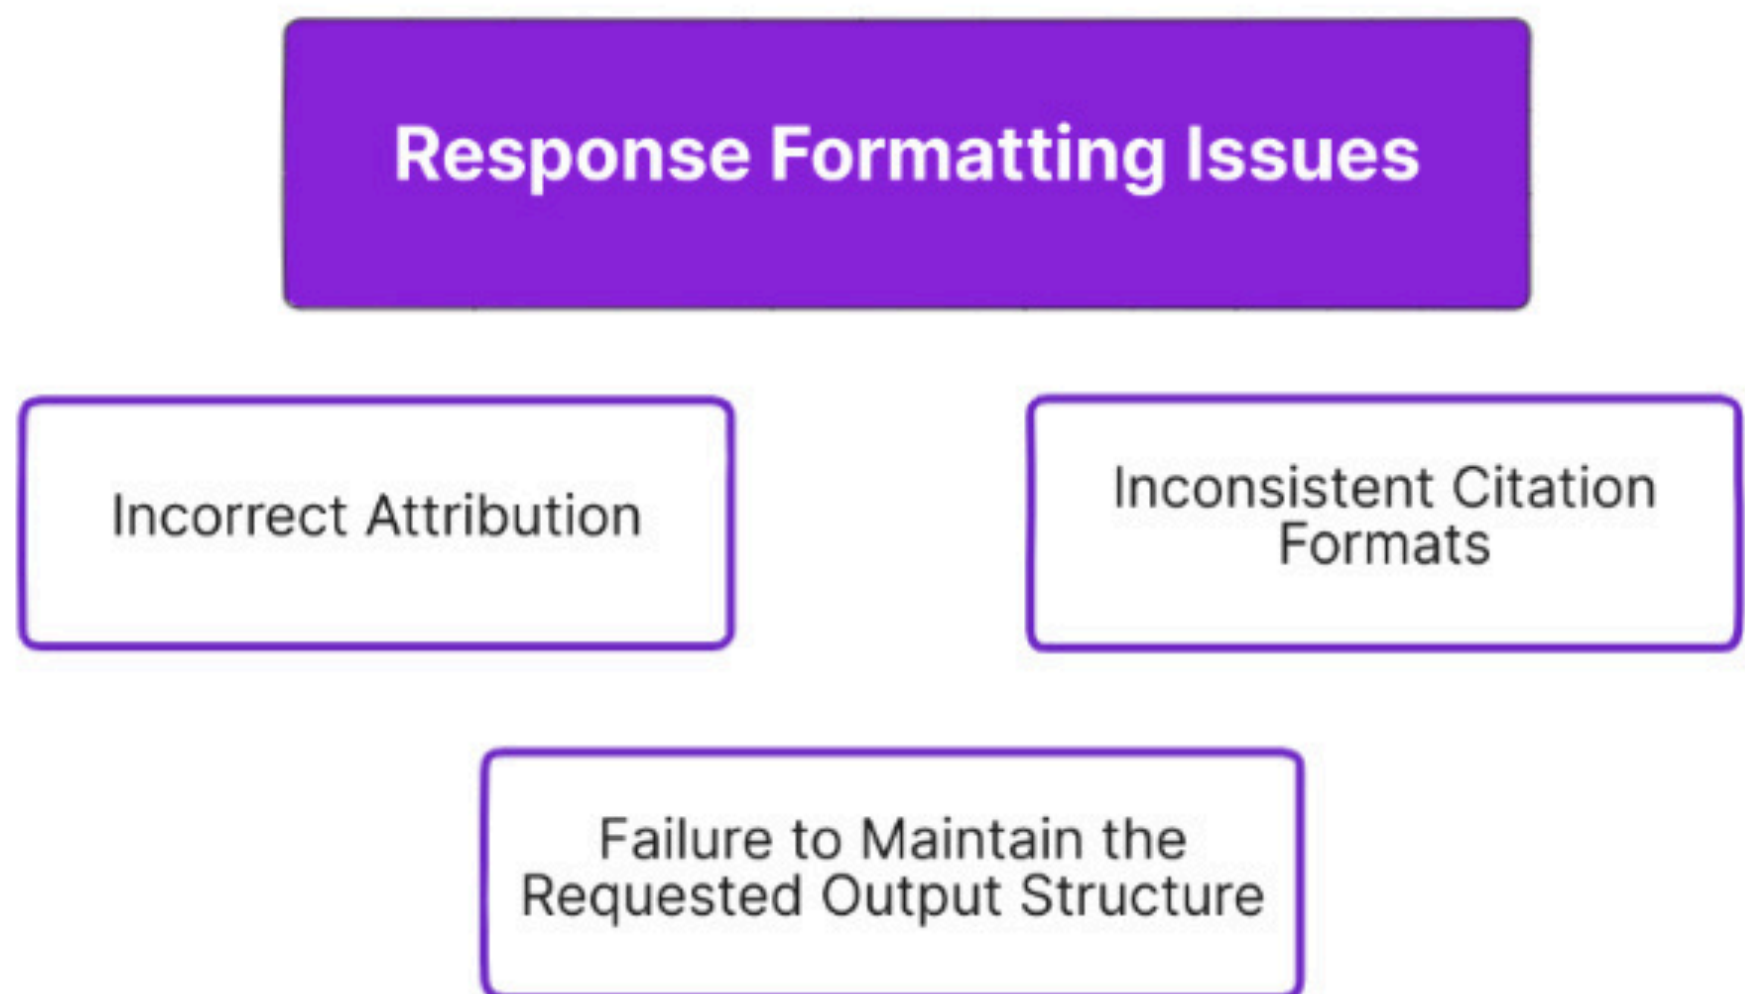
REASONING



Solution:

- Chain-of-thought Prompting
- Multi-step Reasoning Frameworks
- Contradiction Detection Mechanisms

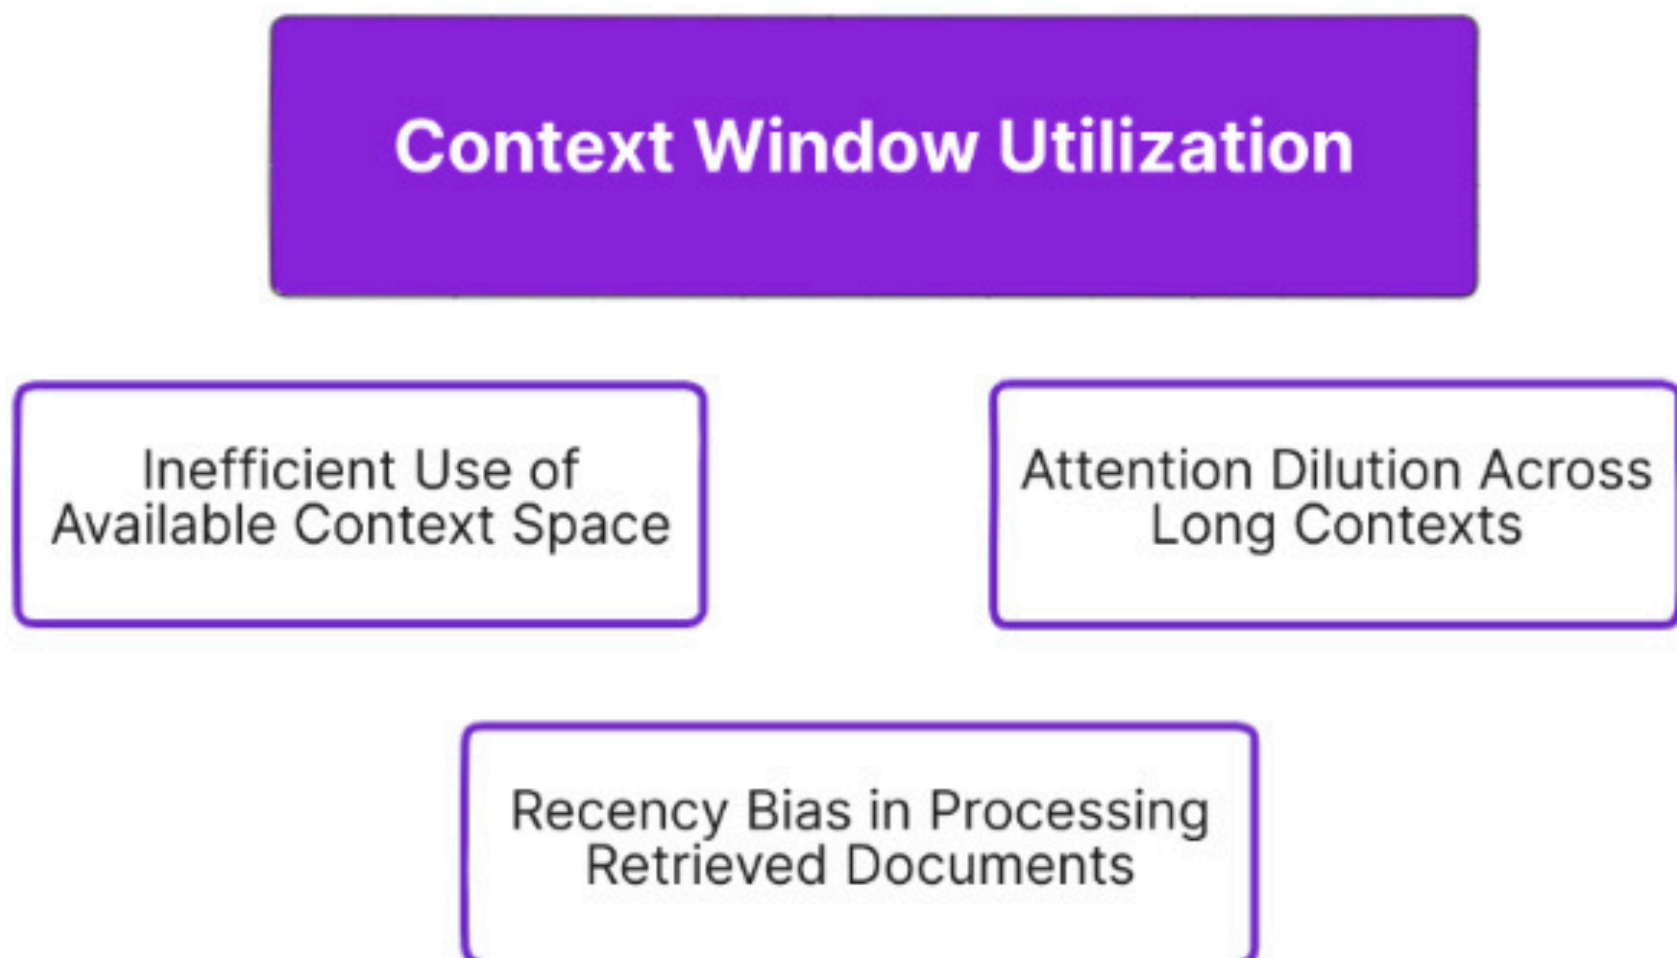
RESPONSE FORMATTING



Solution:

- Enforce structured formatting by using predefined templates
- model with prompt engineering for output formatting
- Automatically checks and corrects attribution, citations, and structure

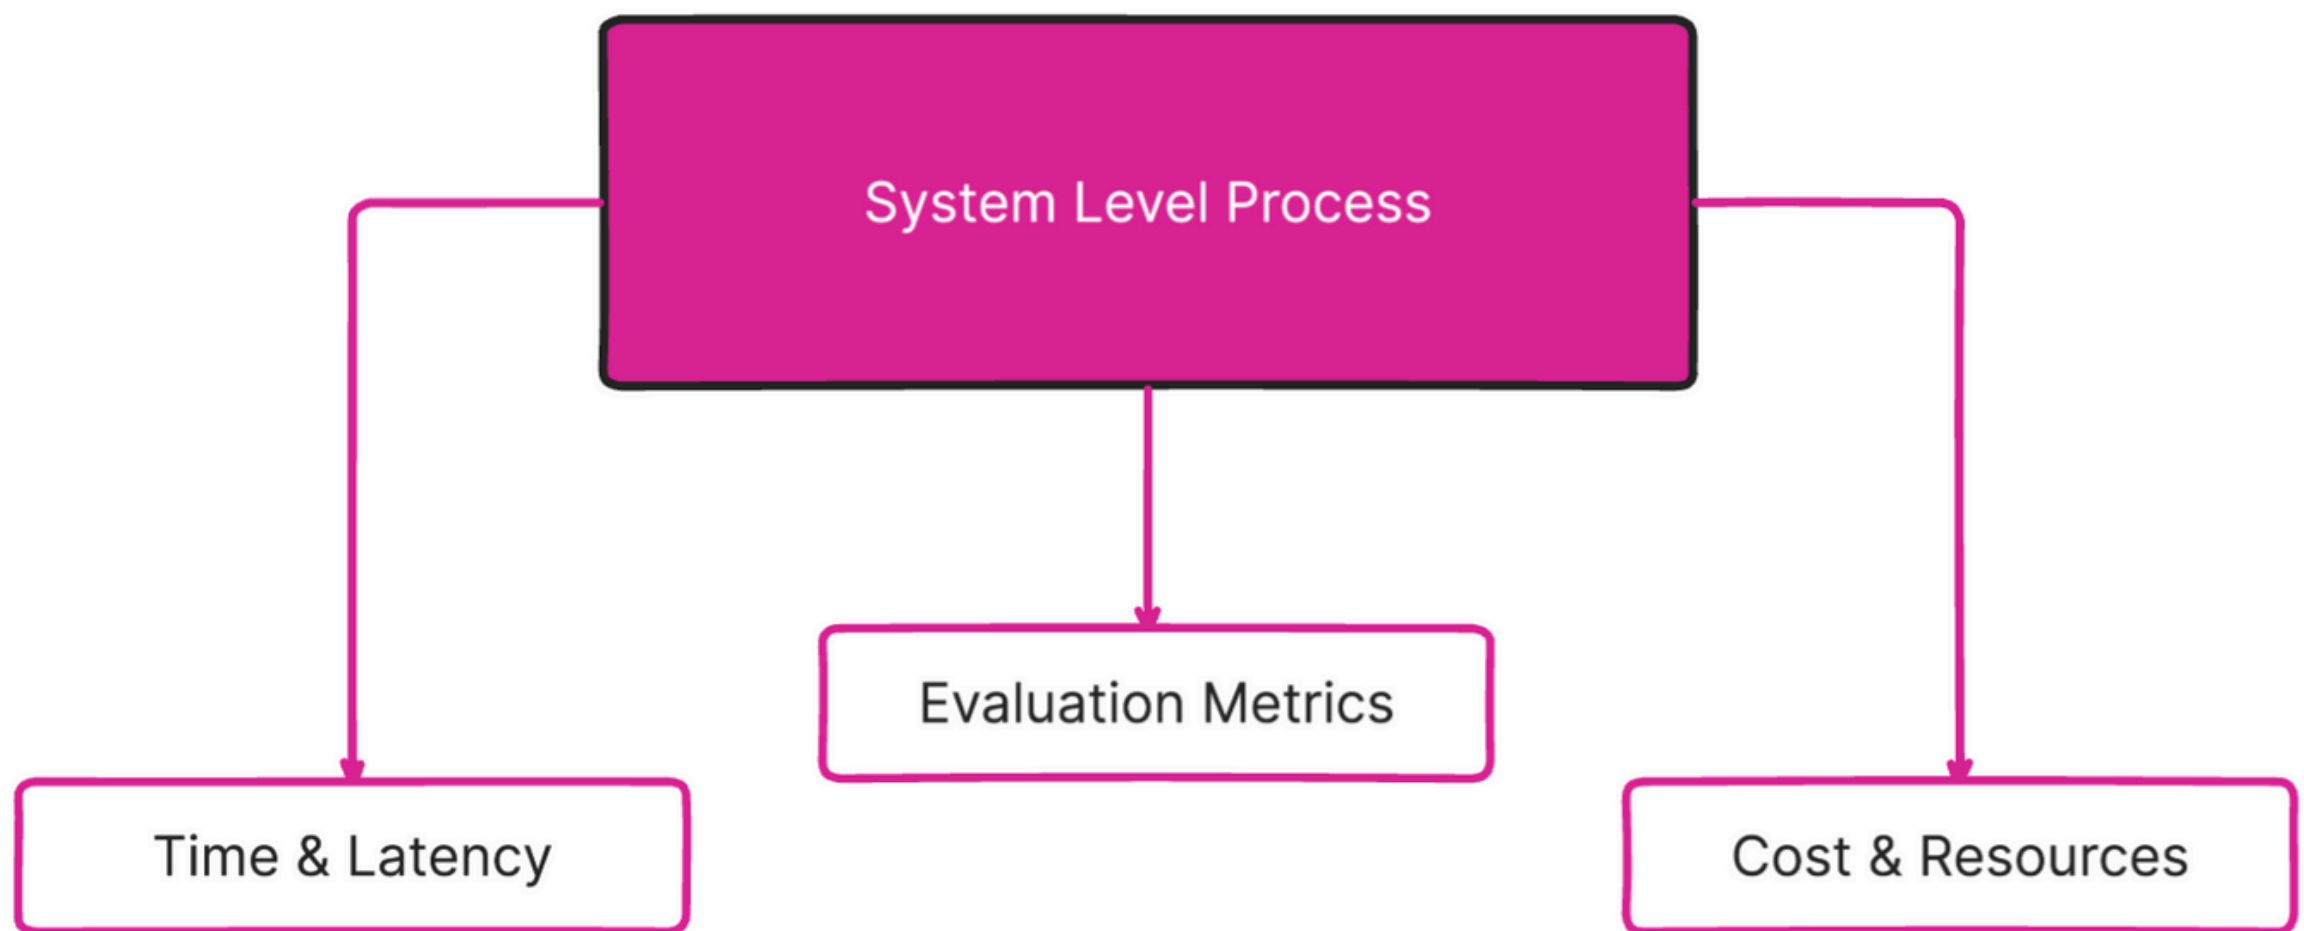
CONTEXT WINDOW



Solution:

- Position key information where the model focuses most
- Maximize value by prioritizing important content and reducing redundancy
- Use structured prompts to direct the model's focus to essential sections.

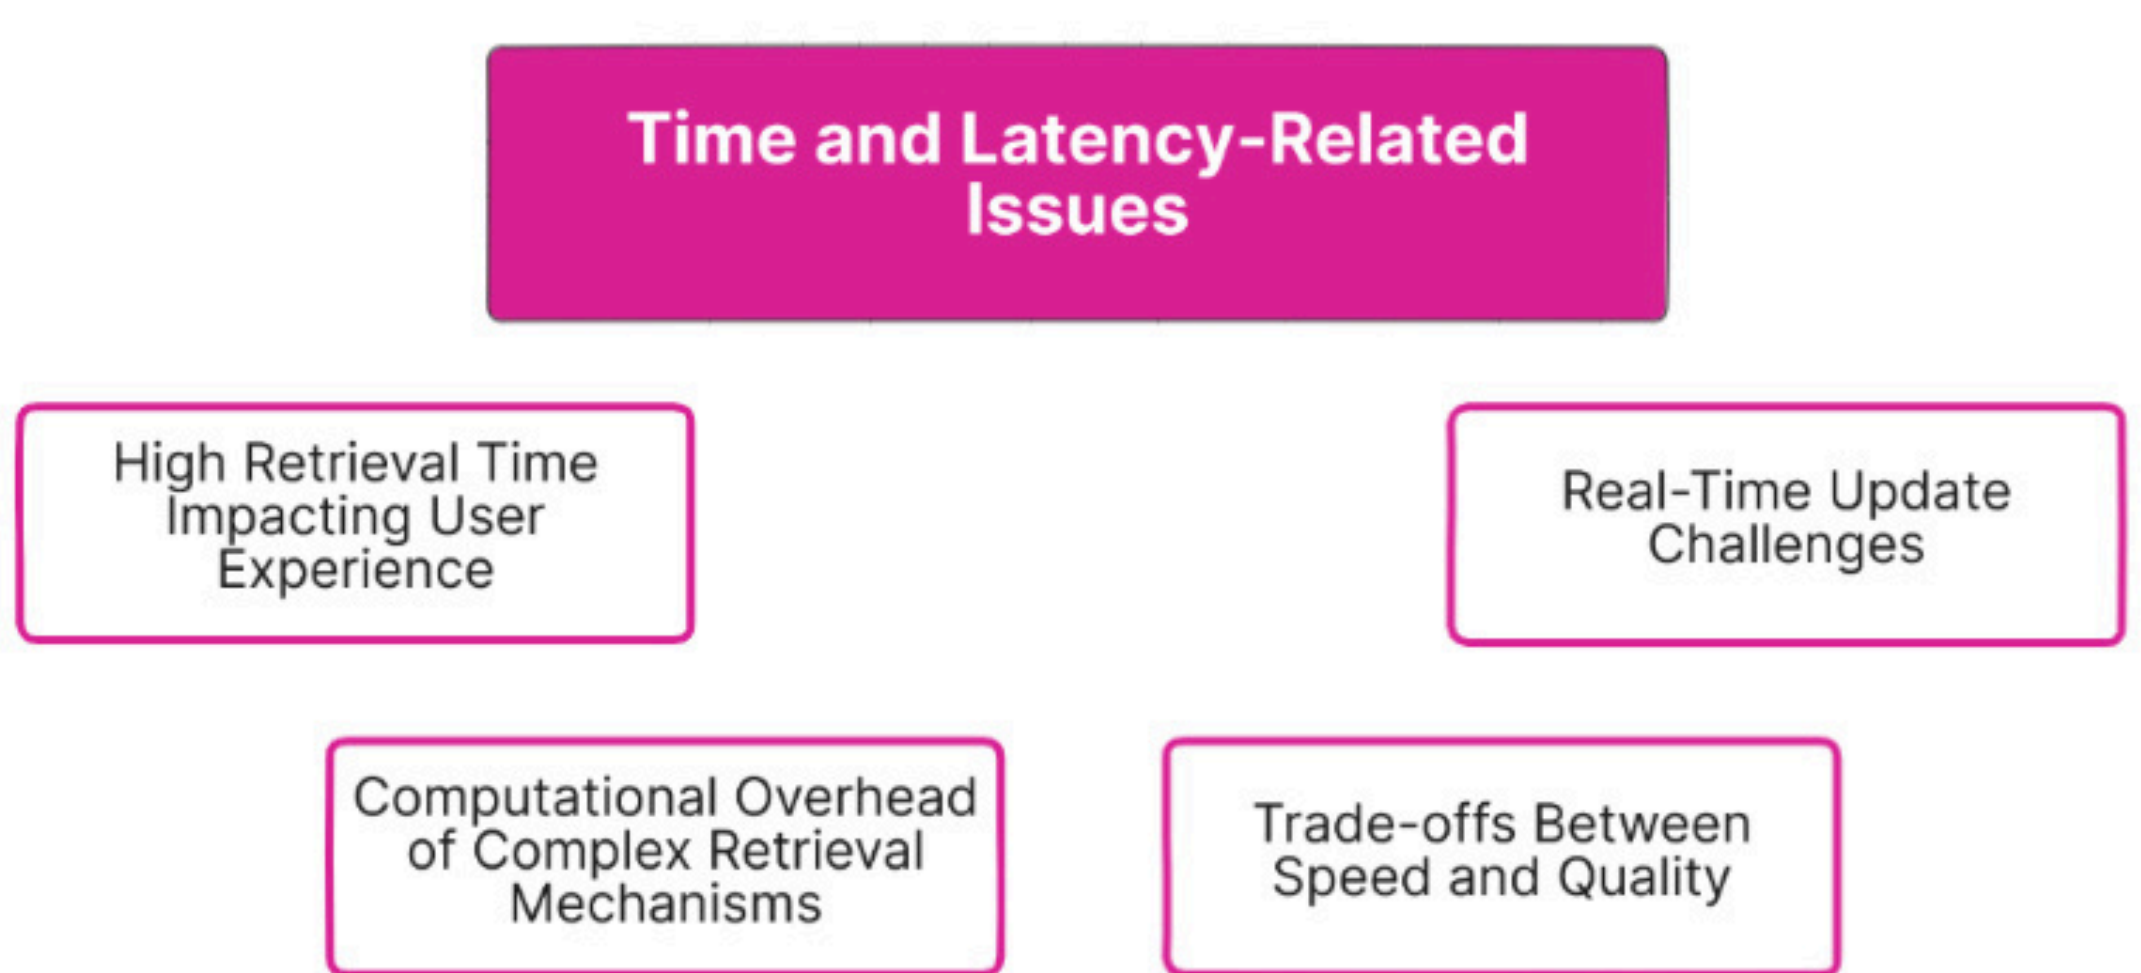
SYSTEM LEVEL FAILURES



Reasons:

- Time & Latency Related Issues
- Evaluation Challenges
- Architectural Limitations
- Cost & Resource Efficiency

TIME & LATENCY RELATED ISSUES



Solution:

- Store common data in memory for faster access.
- Adjust retrieval complexity based on the query.
- Get quick results first, then refine if necessary.
- Refresh knowledge in the background without slowing responses.

EVALUATION CHALLENGES

Evaluation Challenges

Difficulty in Measuring
RAG System Quality
Holistically

Disconnect Between
User Satisfaction and
Technical Metrics

Overemphasis on Retrieval
Metrics at the Expense of
Generation Quality

Solution:

- Assess RAG using retrieval quality, accuracy, coherence, and user engagement.
- Measure user satisfaction via A/B tests, preference modeling, and feedback.
- Test system robustness and grounding by varying retrieval conditions

COST & RESOURCE EFFICIENCY

Cost and Resource Efficiency

Expensive Infrastructure
Requirements

Scaling Challenges for
Enterprise Applications

Storage Constraints for
Large Knowledge Bases


Compute-Intensive
Processing for Large-
Scale Deployment

Solution:

- Use fast, approximate search first, then precise retrieval.
- Compress large models into smaller, efficient versions.
- Utilize methods like BM25/hybrid search to reduce compute/memory.
- Speed up retrieval and lower costs with optimized indexes (ANN, etc.)

READ THE BLOG TO UNDERSTAND HOW TO BUILD AN EFFICIENT RAG SYSTEM

READ MORE



Build a RAG System
That Works!