

COURSE NAME

ADVANCED SQL



GROUP BY, HAVING, ANY, ALL, and EXISTS

Clause	When to Use
<code>GROUP BY</code>	To aggregate data based on one or more columns (e.g., find total salary per department).
<code>HAVING</code>	To filter aggregated results after <code>GROUP BY</code> (e.g., show only departments with total salary > 20,000).
<code>ANY</code>	To compare a value with any value from a subquery (e.g., find employees with a salary higher than at least one employee in another department).
<code>ALL</code>	To compare a value with all values from a subquery (e.g., find employees whose salary is higher than everyone in another department).
<code>EXISTS</code>	To check if a subquery returns any rows (e.g., find departments that have at least one employee).

EmployeeID	Name	Department	Salary
2	Bob	HR	7000
3	Charlie	HR	6000
1	Alice	HR	5000
6	Frank	IT	9000
4	David	IT	8000
5	Eve	IT	7500
8	Hank	Finance	7200
7	Grace	Finance	6500

1. GROUP BY with Aggregation (SUM(), AVG(), COUNT())

```
SELECT Department, COUNT(*) AS EmployeeCount, SUM(Salary)
AS TotalSalary, AVG(Salary) AS AvgSalary
FROM Employees
GROUP BY Department;
```

Result

Department	EmployeeCount	TotalSalary	AvgSalary
HR	3	18500	6166.67
IT	3	24500	8166.67
Finance	2	13700	6850.00

Explanation:

- Groups employees by `Department` and calculates:
 - `COUNT(*)` → Number of employees per department.
 - `SUM(Salary)` → Total salary per department.
 - `AVG(Salary)` → Average salary per department.

2. HAVING (Filter After GROUP BY)

Find departments where the **total salary** exceeds 20,000.

```
SELECT Department, SUM(Salary) AS TotalSalary
FROM Employees
GROUP BY Department
HAVING SUM(Salary) > 20000;
```

Result

Department	TotalSalary
HR	18500
IT	24500

Explanation:

- **HAVING** is used to filter aggregated values (`SUM(Salary) > 20000`).
- Unlike **WHERE** , which filters individual rows, **HAVING** works on grouped data.

3. ANY (Compare Against Any Value in a Subquery)

Find employees who earn **more than at least one employee in the HR department**.

```
SELECT Name, Salary
FROM Employees
WHERE Salary > ANY (SELECT Salary FROM Employees WHERE
Department = 'HR');
```

Result

Name	Salary
Bob	7000
Charlie	6000
David	8000
Eve	7500
Frank	9000
Hank	7200
Grace	6500

Explanation:

- The subquery `(SELECT Salary FROM Employees WHERE Department = 'HR')` returns `{5000, 5500, 7000}`.
- `ANY` means at least one comparison must be true, so it finds employees earning more than any salary in HR (i.e., greater than `5000`, `5500`, or `7000`).

4. ALL (Compare Against All Values in a Subquery)

Find employees who earn **more than every employee in the Finance department**.

```
SELECT Name, Salary
FROM Employees
WHERE Salary > ALL (SELECT Salary FROM Employees WHERE
Department = 'Finance');
```

Result

Name	Salary
Bob	7000
Charlie	6000
David	8000
Eve	7500
Frank	9000

Explanation:

- The subquery `(SELECT Salary FROM Employees WHERE Department = 'Finance')` returns `{6500, 7200}`.
- `ALL` means the condition must be true for every row.
- Employees must have a salary greater than 7200.

5. EXISTS (Check If Any Matching Row Exists)

Find departments that **have at least one employee**.

```
SELECT DISTINCT Department
FROM Employees e
WHERE EXISTS (SELECT 1 FROM Employees WHERE e.Department =
Employees.Department);
```

Result

Department
HR
IT
Finance

Explanation:

- `EXISTS` checks if the subquery returns any row.
- Since every department has employees, all departments are listed.

6. Combining GROUP BY, HAVING, and EXISTS

Find departments where the **average salary is above 6000** and **there are employees in that department**.

```
SELECT Department, AVG(Salary) AS AvgSalary
FROM Employees
WHERE EXISTS (SELECT 1 FROM Employees e WHERE e.Department
= Employees.Department)
GROUP BY Department
HAVING AVG(Salary) > 6000;
```

Department	AvgSalary
HR	6166.67
IT	8166.67
Finance	6850.00

Explanation:

- **EXISTS** ensures the department has employees.
- **GROUP BY** groups employees by department.
- **HAVING** filters out departments where $AVG(Salary) \leq 6000$.

Conclusion

Clause	Purpose
GROUP BY	Groups rows and applies aggregate functions (SUM() , AVG() , etc.).
HAVING	Filters grouped results based on aggregate conditions.
ANY	Checks if any value in a subquery meets a condition.
ALL	Checks if all values in a subquery meet a condition.
EXISTS	Returns TRUE if the subquery returns any rows.

Correlated and Nested Queries

Query Type	When to Use	Example Use Case
Nested Queries (Subqueries)	When you need a query inside another query that runs independently	Finding employees who earn more than the average salary
Correlated Subqueries	When the inner query depends on the outer query and runs for each row	Finding employees who earn the highest salary in their department

1. Nested Queries (Subqueries)

Example 1: Find Employees Who Earn More Than the Average Salary

```
SELECT Name, Salary
FROM Employees
WHERE Salary > (SELECT AVG(Salary) FROM Employees);
```

Result

Name	Salary
Bob	7000
David	8000
Eve	7500
Frank	9000
Hank	7200

How It Works:

- The subquery `(SELECT AVG(Salary) FROM Employees)` calculates the average salary.
- The main query then selects employees whose salary is greater than this value.

2. Correlated Subqueries

Example 2: Find Employees Who Earn the Highest Salary in Their Department

```
SELECT Name, Salary, DepartmentID
FROM Employees e1
WHERE Salary = (SELECT MAX(Salary) FROM Employees e2 WHERE
e1.DepartmentID = e2.DepartmentID);
```

Result

Name	Salary	DepartmentID
Bob	7000	1
Frank	9000	2
Hank	7200	3

How It Works:

- The **inner query** `(SELECT MAX(Salary) FROM Employees e2 WHERE e1.DepartmentID = e2.DepartmentID)`
 - Runs for each row in the outer query.
 - Finds the maximum salary in the same department as the current row.
- The outer query then selects only those employees whose salary matches the max salary.

SELF JOIN

Find Pairs of Employees in the Same Department

```
SELECT e1.Name AS Employee1, e2.Name AS Employee2,  
e1.DepartmentID  
FROM Employees e1  
JOIN Employees e2 ON e1.DepartmentID = e2.DepartmentID AND  
e1.EmployeeID <> e2.EmployeeID;
```

Result

Employee1	Employee2	DepartmentID
Alice	Bob	1
Alice	Charlie	1
Bob	Alice	1
Bob	Charlie	1

How It Works:

- Joins `Employees` to itself to find pairs in the same department.

PARTITION BY

The **PARTITION BY** clause is used in **window functions** to divide the result set into partitions and perform calculations within each partition **without collapsing rows**. Use it when you need **aggregations over subsets of data** while retaining row-level detail.

Common Use Cases

1. **Ranking rows within groups** (e.g., finding the top N sales per region)
2. **Running totals or cumulative sums** (e.g., calculating a running balance)
3. **Comparing current and previous rows** (e.g., difference from the last transaction)
4. **Aggregations while keeping details** (e.g., average salary per department)

Key Difference: PARTITION BY vs GROUP BY

Feature	PARTITION BY	GROUP BY
Keeps individual rows?	✓ Yes	✗ No (Collapses rows)
Used with	Window Functions	Aggregate Functions
Example Functions	RANK(), SUM() OVER, LAG()	SUM(), COUNT(), AVG()
Example Use	Running totals, rankings, previous row values	Total sales per region

Examples

EmployeeID	Name	Department	Salary
2	Bob	HR	7000
3	Charlie	HR	6000
1	Alice	HR	5000
6	Frank	IT	9000
4	David	IT	8000
5	Eve	IT	7500
8	Hank	Finance	7200
7	Grace	Finance	6500

Example 1: Ranking Employees by Salary Within Department

```
SELECT EmployeeID, Name, Department, Salary,  
RANK() OVER (PARTITION BY Department ORDER BY Salary DESC)  
AS Rank  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	Rank
2	Bob	HR	7000	1
3	Charlie	HR	6000	2
1	Alice	HR	5000	3
6	Frank	IT	9000	1
4	David	IT	8000	2
5	Eve	IT	7500	3
8	Hank	Finance	7200	1
7	Grace	Finance	6500	2

Example 2: Running Total of Salary in Each Department

```
SELECT EmployeeID, Name, Department, Salary,  
SUM(Salary) OVER (PARTITION BY Department ORDER BY Salary)  
AS RunningTotal  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	RunningTotal
1	Alice	HR	5000	5000
3	Charlie	HR	6000	11000
2	Bob	HR	7000	18000
5	Eve	IT	7500	7500
4	David	IT	8000	15500
6	Frank	IT	9000	24500
7	Grace	Finance	6500	6500
8	Hank	Finance	7200	13700

Example 3: Previous Employee Salary in Each Department

```
SELECT EmployeeID, Name, Department, Salary,  
LAG(Salary) OVER (PARTITION BY Department ORDER BY Salary)  
AS PreviousSalary  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	PreviousSalary
1	Alice	HR	5000	NULL
3	Charlie	HR	6000	5000
2	Bob	HR	7000	6000
5	Eve	IT	7500	NULL
4	David	IT	8000	7500
6	Frank	IT	9000	8000
7	Grace	Finance	6500	NULL
8	Hank	Finance	7200	6500

Example 4: Average Salary per Department

```
SELECT EmployeeID, Name, Department, Salary,  
AVG(Salary) OVER (PARTITION BY Department) AS AvgSalary  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	AvgSalary
1	Alice	HR	5000	6000
2	Bob	HR	7000	6000
3	Charlie	HR	6000	6000
4	David	IT	8000	8166.67
5	Eve	IT	7500	8166.67
6	Frank	IT	9000	8166.67
7	Grace	Finance	6500	6850
8	Hank	Finance	7200	6850

Summary

- **RANK()** assigns **ranking within partitions** (departments).
- **SUM() OVER** calculates **cumulative totals** within partitions.
- **LAG()** retrieves **previous row's value** within partitions.
- **AVG() OVER** computes **aggregates while keeping row details**.

RANK(), DENSE_RANK(), and ROW_NUMBER()

These three **window functions** are used for **assigning unique row numbers** based on a given ordering within partitions (groups of data).

Function	When to Use
RANK()	When you need rankings with gaps if there are ties.
DENSE_RANK()	When you need rankings without gaps if there are ties.
ROW_NUMBER()	When you need a unique row number for each row , even if values are the same.

1. **RANK()** Example (Gaps in Ranking)

```
SELECT EmployeeID, Name, Department, Salary,  
RANK() OVER (PARTITION BY Department ORDER BY Salary DESC)  
AS Rank  
FROM Employees;
```

Result (RANK() has gaps)

EmployeeID	Name	Department	Salary	Rank
2	Bob	HR	7000	1
3	Charlie	HR	6000	2
9	Ivy	HR	6000	2
1	Alice	HR	5000	4
6	Frank	IT	9000	1
4	David	IT	8000	2
5	Eve	IT	7500	3
8	Hank	Finance	7200	1
7	Grace	Finance	6500	2

Note: Since **Charlie** and **Ivy** have the same salary, they share rank 2, and the next rank skips to 4.

2. **DENSE_RANK()** Example (No Gaps in Ranking)

```
SELECT EmployeeID, Name, Department, Salary,  
DENSE_RANK() OVER (PARTITION BY Department ORDER BY Salary  
DESC) AS DenseRank  
FROM Employees;
```

Result (`DENSE_RANK()` removes gaps)

EmployeeID	Name	Department	Salary	DenseRank
2	Bob	HR	7000	1
3	Charlie	HR	6000	2
9	Ivy	HR	6000	2
1	Alice	HR	5000	3
6	Frank	IT	9000	1
4	David	IT	8000	2
5	Eve	IT	7500	3
8	Hank	Finance	7200	1
7	Grace	Finance	6500	2

Difference: `DENSE_RANK()` ensures there are no gaps—after 2, the next rank is 3, not 4.

3. `ROW_NUMBER()` Example (Unique Row Number)

```
SELECT EmployeeID, Name, Department, Salary,  
ROW_NUMBER() OVER (PARTITION BY Department ORDER BY Salary  
DESC) AS RowNum  
FROM Employees;
```

Result (ROW_NUMBER() assigns unique numbers)

EmployeeID	Name	Department	Salary	RowNum
2	Bob	HR	7000	1
3	Charlie	HR	6000	2
9	Ivy	HR	6000	3
1	Alice	HR	5000	4
6	Frank	IT	9000	1
4	David	IT	8000	2
5	Eve	IT	7500	3
8	Hank	Finance	7200	1
7	Grace	Finance	6500	2

Comparison Table

Function	Handles Ties?	Gaps in Ranking?	Unique Row for Each Record?
RANK()	Yes	Yes	No
DENSE_RANK()	Yes	No (Ranks stay sequential)	No
ROW_NUMBER()	No	No (Strictly sequential)	Yes

When to Use Which?

Use Case	Best Function
Need rankings with gaps	RANK()
Need rankings without gaps	DENSE_RANK()
Need strict numbering with no ties	ROW_NUMBER()
Need to pick the top N per group	ROW_NUMBER()
Need consistent ranking even with ties	DENSE_RANK()

Example: Get Top 2 Highest Salaries per Department

```
SELECT EmployeeID, Name, Department, Salary, Rank
FROM (
    SELECT EmployeeID, Name, Department, Salary,
    RANK() OVER
    (PARTITION BY Department ORDER BY Salary DESC) AS Rank
FROM Employees)
RankedEmployees WHERE Rank <= 2;
```

Result

EmployeeID	Name	Department	Salary	Rank
2	Bob	HR	7000	1
3	Charlie	HR	6000	2
9	Ivy	HR	6000	2
6	Frank	IT	9000	1
4	David	IT	8000	2
8	Hank	Finance	7200	1
7	Grace	Finance	6500	2

LEAD() Function in SQL - Example

The `LEAD()` function is used to access data from **the next row** within the same result set **without using a self-join**. It is useful for **comparing values between rows**.

Example Use Cases

- Finding the **next employee's salary** in the same department.
- Calculating **difference between current and next month's sales**.
- Checking **the next event in a schedule**.

1. Basic `LEAD()` Example

```
SELECT EmployeeID, Name, Department, Salary,  
LEAD(Salary) OVER (PARTITION BY Department ORDER BY Salary  
DESC) AS NextSalary  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	NextSalary
2	Bob	HR	7000	6000
3	Charlie	HR	6000	5000
1	Alice	HR	5000	NULL
6	Frank	IT	9000	8000
4	David	IT	8000	7500
5	Eve	IT	7500	NULL
8	Hank	Finance	7200	6500
7	Grace	Finance	6500	NULL

Explanation:

- `LEAD(Salary)` returns the **next salary in the same department** based on `ORDER BY Salary DESC`.
- For the last employee in each department, `LEAD()` returns `NULL` (no next row available).

2. Using **LEAD()** to Find Salary Difference

```
SELECT EmployeeID, Name, Department, Salary,  
LEAD(Salary) OVER (PARTITION BY Department ORDER BY Salary  
DESC) AS NextSalary,  
Salary - LEAD(Salary) OVER (PARTITION BY Department  
ORDER BY Salary DESC) AS SalaryDifference  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	NextSalary	SalaryDifference
2	Bob	HR	7000	6000	1000
3	Charlie	HR	6000	5000	1000
1	Alice	HR	5000	NULL	NULL
6	Frank	IT	9000	8000	1000
4	David	IT	8000	7500	500
5	Eve	IT	7500	NULL	NULL
8	Hank	Finance	7200	6500	700
7	Grace	Finance	6500	NULL	NULL

Explanation:

- The column `SalaryDifference` calculates the difference between current and next salary.
- If there is no next salary (`NULL`), then `SalaryDifference` is also `NULL` .

3. Using **LEAD()** with a Custom Offset

You can specify how many rows ahead you want to look at.

```
SELECT EmployeeID, Name, Department, Salary,  
LEAD(Salary, 2) OVER (PARTITION BY Department ORDER BY  
Salary DESC) AS SalaryAfterTwo  
FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	SalaryAfterTwo
2	Bob	HR	7000	5000
3	Charlie	HR	6000	NULL
1	Alice	HR	5000	NULL
6	Frank	IT	9000	7500
4	David	IT	8000	NULL
5	Eve	IT	7500	NULL
8	Hank	Finance	7200	NULL
7	Grace	Finance	6500	NULL

Explanation:

- `LEAD(Salary, 2)` looks two rows ahead instead of just one.
- If there aren't two rows ahead, it returns `NULL`.

4. Handling NULL Values with **LEAD()** (Default Value)

By default, **LEAD()** returns `NULL` if there's no next row. You can provide a **default value** instead.

```
SELECT EmployeeID, Name, Department, Salary,  
LEAD(Salary, 1, 0) OVER (PARTITION BY Department ORDER BY  
Salary DESC) AS NextSalary FROM Employees;
```

Result

EmployeeID	Name	Department	Salary	NextSalary
2	Bob	HR	7000	6000
3	Charlie	HR	6000	5000
1	Alice	HR	5000	0
6	Frank	IT	9000	8000
4	David	IT	8000	7500
5	Eve	IT	7500	0
8	Hank	Finance	7200	6500
7	Grace	Finance	6500	0

Explanation:

- `LEAD(Salary, 1, 0)` means if there's no next salary, it returns 0 instead of `NULL`.

When to Use `LEAD()` ?

Use Case	Why Use <code>LEAD()</code> ?
Comparing current vs next row values	E.g., Checking the salary difference between employees.
Finding next events, orders, or transactions	E.g., Checking what the next scheduled appointment is.
Performing trend analysis	E.g., Comparing sales of the current month vs the next month.
Avoiding self-joins for next row comparisons	E.g., Instead of using <code>JOIN</code> on the same table, <code>LEAD()</code> simplifies the query.

Summary

- `LEAD()` looks ahead to the next row.
- It helps in comparing values between rows without using a self-join.
- You can customize the offset (`LEAD(value, offset)`).
- You can set a default value to replace `NULL` (`LEAD(value, offset, default_value)`).

Exercise

Here's a **comprehensive list of 50 SQL exercises** from platforms like **LeetCode**, **DataLemur**, and **SQLZoo**, covering **basic to advanced** SQL concepts. Each exercise includes the **problem statement, solution, and explanation**.

List of Topics Covered

- ✓ **Basic Queries** (`SELECT`, `WHERE`, `ORDER BY`)
 - ✓ **Aggregations** (`GROUP BY`, `HAVING`)
 - ✓ **Joins** (`INNER JOIN`, `LEFT JOIN`, `RIGHT JOIN`, `FULL JOIN`, `SELF JOIN`, `CROSS JOIN`)
 - ✓ **Subqueries** (`Nested`, `Correlated`)
 - ✓ **Window Functions** (`RANK()`, `DENSE_RANK()`, `ROW_NUMBER()`, `LEAD()`, `LAG()`)
 - ✓ **Set Operations** (`UNION`, `INTERSECT`, `EXCEPT`)
 - ✓ **Conditional Statements** (`CASE WHEN`)
 - ✓ **Existence Checks** (`EXISTS`, `NOT EXISTS`)
 - ✓ **Filtering** (`ANY`, `ALL`, `LIKE`, `IN`)
-

Basic SQL Queries

1 Retrieve All Employees' Names and Salaries

Problem: Select all employees' names and salaries from the `Employees` table.

```
SELECT Name, Salary FROM Employees;
```

-
- ♦ **Explanation:** Basic `SELECT` statement to retrieve data.
-

2 Find Employees Earning More Than \$50,000

```
SELECT Name, Salary
FROM Employees
WHERE Salary > 50000;
```

- ♦ **Explanation:** Uses `WHERE` to filter salaries greater than 50000.
-

3 Order Employees by Salary Descending

```
SELECT Name, Salary
FROM Employees
ORDER BY Salary DESC;
```

- ♦ **Explanation:** `ORDER BY DESC` sorts from highest to lowest salary.
-



Aggregation & Grouping

4 Count Employees in Each Department

```
SELECT Department, COUNT(*) AS EmployeeCount
FROM Employees
GROUP BY Department;
```


-
- ♦ **Explanation:** Uses `COUNT(*)` with `GROUP BY` to count employees per department.
-

5 Find Departments with More Than 10 Employees

```
SELECT Department, COUNT(*) AS EmployeeCount
FROM Employees
GROUP BY Department
HAVING COUNT(*) > 10;
```

- ♦ **Explanation:**
 - ✓ `HAVING` filters groups after aggregation (unlike `WHERE`, which filters before grouping).
-

6 Find the Maximum Salary in Each Department

```
SELECT Department, MAX(Salary) AS MaxSalary
FROM Employees
GROUP BY Department;
```

- ♦ **Explanation:** Uses `MAX()` to find the highest salary per department.
-

SQL Joins

7 Retrieve Employees with Their Department Names

```
SELECT e.Name, e.Salary, d.DepartmentName
FROM Employees e
INNER JOIN Departments d
ON e.DepartmentID = d.DepartmentID;
```

- ♦ **Explanation:** **INNER JOIN** matches employees with departments.
-

8 List All Employees, Even Those Without a Department (**LEFT JOIN**)

```
SELECT e.Name, e.Salary, d.DepartmentName
FROM Employees e
LEFT JOIN Departments d
ON e.DepartmentID = d.DepartmentID;
```

- ♦ **Explanation:**
 - ✓ **LEFT JOIN** ensures **all employees** appear, even if they have no department.
-

9 Find Departments Without Employees (**RIGHT JOIN**)

```
SELECT d.DepartmentName
FROM Employees e
RIGHT JOIN Departments d
ON e.DepartmentID = d.DepartmentID
WHERE e.EmployeeID IS NULL;
```

♦ **Explanation:**

- ✓ **RIGHT JOIN** ensures all departments appear.
 - ✓ **WHERE e.EmployeeID IS NULL** filters **departments with no employees**.
-

Subqueries

10 Find Employees Earning More Than the Average Salary

```
SELECT Name, Salary
FROM Employees
WHERE Salary > (SELECT AVG(Salary) FROM Employees);
```

- ♦ **Explanation:** Uses a **subquery** to compute average salary.
-

11 Find the Department with the Highest Salary

```
SELECT DepartmentID
FROM Employees
WHERE Salary = (SELECT MAX(Salary) FROM Employees);
```

- ♦ **Explanation: Nested Query** retrieves max salary and finds the department.
-

Window Functions

12 Rank Employees by Salary in Each Department

```
SELECT Name, Salary, DepartmentID,  
       RANK() OVER (PARTITION BY DepartmentID ORDER BY Salary DESC) AS  
SalaryRank  
FROM Employees;
```

♦ **Explanation:**

- ✓ **RANK()** assigns ranks within **each department** based on **Salary DESC**.

13 Assign a Unique Row Number to Employees Per Department

```
SELECT Name, Salary, DepartmentID,  
       ROW_NUMBER() OVER (PARTITION BY DepartmentID ORDER BY Salary DESC)  
AS RowNum  
FROM Employees;
```

♦ **Explanation:**

- ✓ **ROW_NUMBER()** assigns **unique numbers** per department.

Set Operations

14 Find Employees Who Are Also in the Managers Table (**INTERSECT**)

```
SELECT Name FROM Employees
INTERSECT
SELECT Name FROM Managers;
```

♦ **Explanation:**

- ✓ **INTERSECT** returns names common to both tables.
-

15 Find Employees Who Are Not Managers (**EXCEPT**)

```
SELECT Name FROM Employees
EXCEPT
SELECT Name FROM Managers;
```

♦ **Explanation:**

- ✓ **EXCEPT** returns employees **who are not in Managers**.
-

EXISTS, ANY, ALL

16 Find Employees in Departments with More Than 5 Employees

```
SELECT Name, DepartmentID
FROM Employees e
WHERE EXISTS (SELECT 1 FROM Employees e2 WHERE e.DepartmentID =
e2.DepartmentID GROUP BY e2.DepartmentID HAVING COUNT(*) > 5);
```

♦ **Explanation:**

✓ **EXISTS** checks if any department has **more than 5 employees**.

Find Employees Who Earn More Than Every Employee in HR (ALL)

```
SELECT Name, Salary
FROM Employees
WHERE Salary > ALL (SELECT Salary FROM Employees WHERE DepartmentID = 1);
```

♦ **Explanation:**

✓ Compares salary with **all salaries** in HR.

Advanced Queries

Find the Second Highest Salary

```
SELECT DISTINCT Salary
FROM Employees
ORDER BY Salary DESC
LIMIT 1 OFFSET 1;
```

♦ **Explanation:**

✓ **OFFSET 1** skips the highest salary and fetches the second highest.

19 Find Employees Who Joined Before Their Manager

```
SELECT e.Name
FROM Employees e
JOIN Employees m ON e.ManagerID = m.EmployeeID
WHERE e.JoinDate < m.JoinDate;
```

♦ **Explanation:**

- ✓ Self-join (`JOIN Employees m ON e.ManagerID = m.EmployeeID`).

20 Retrieve Employee Salary Difference from the Previous Month (`LAG()`)

```
SELECT EmployeeID, Name, Salary,
       Salary - LAG(Salary) OVER (PARTITION BY EmployeeID ORDER BY Month)
AS SalaryChange
FROM Salaries;
```

♦ **Explanation:**

- ✓ `LAG()` fetches the **previous month's salary** for comparison.